
Classifying galaxies

PHYS 246 class 7

<https://lkwagner.github.io/IntroductionToComputationalPhysics/intro.html>

Announcements/notes

- 'Particle physics' is due tonight.
- I updated the 'Classifying galaxies' notebook. **you must turn in the new version with spring 2025 on the top.** (it's better!)
-

```
from google.colab import drive  
drive.mount('/content/drive')  
!cp /content/drive/MyDrive/Colab\ Notebooks/Dynamics.ipynb ./  
!jupyter nbconvert --to HTML "Dynamics.ipynb"
```

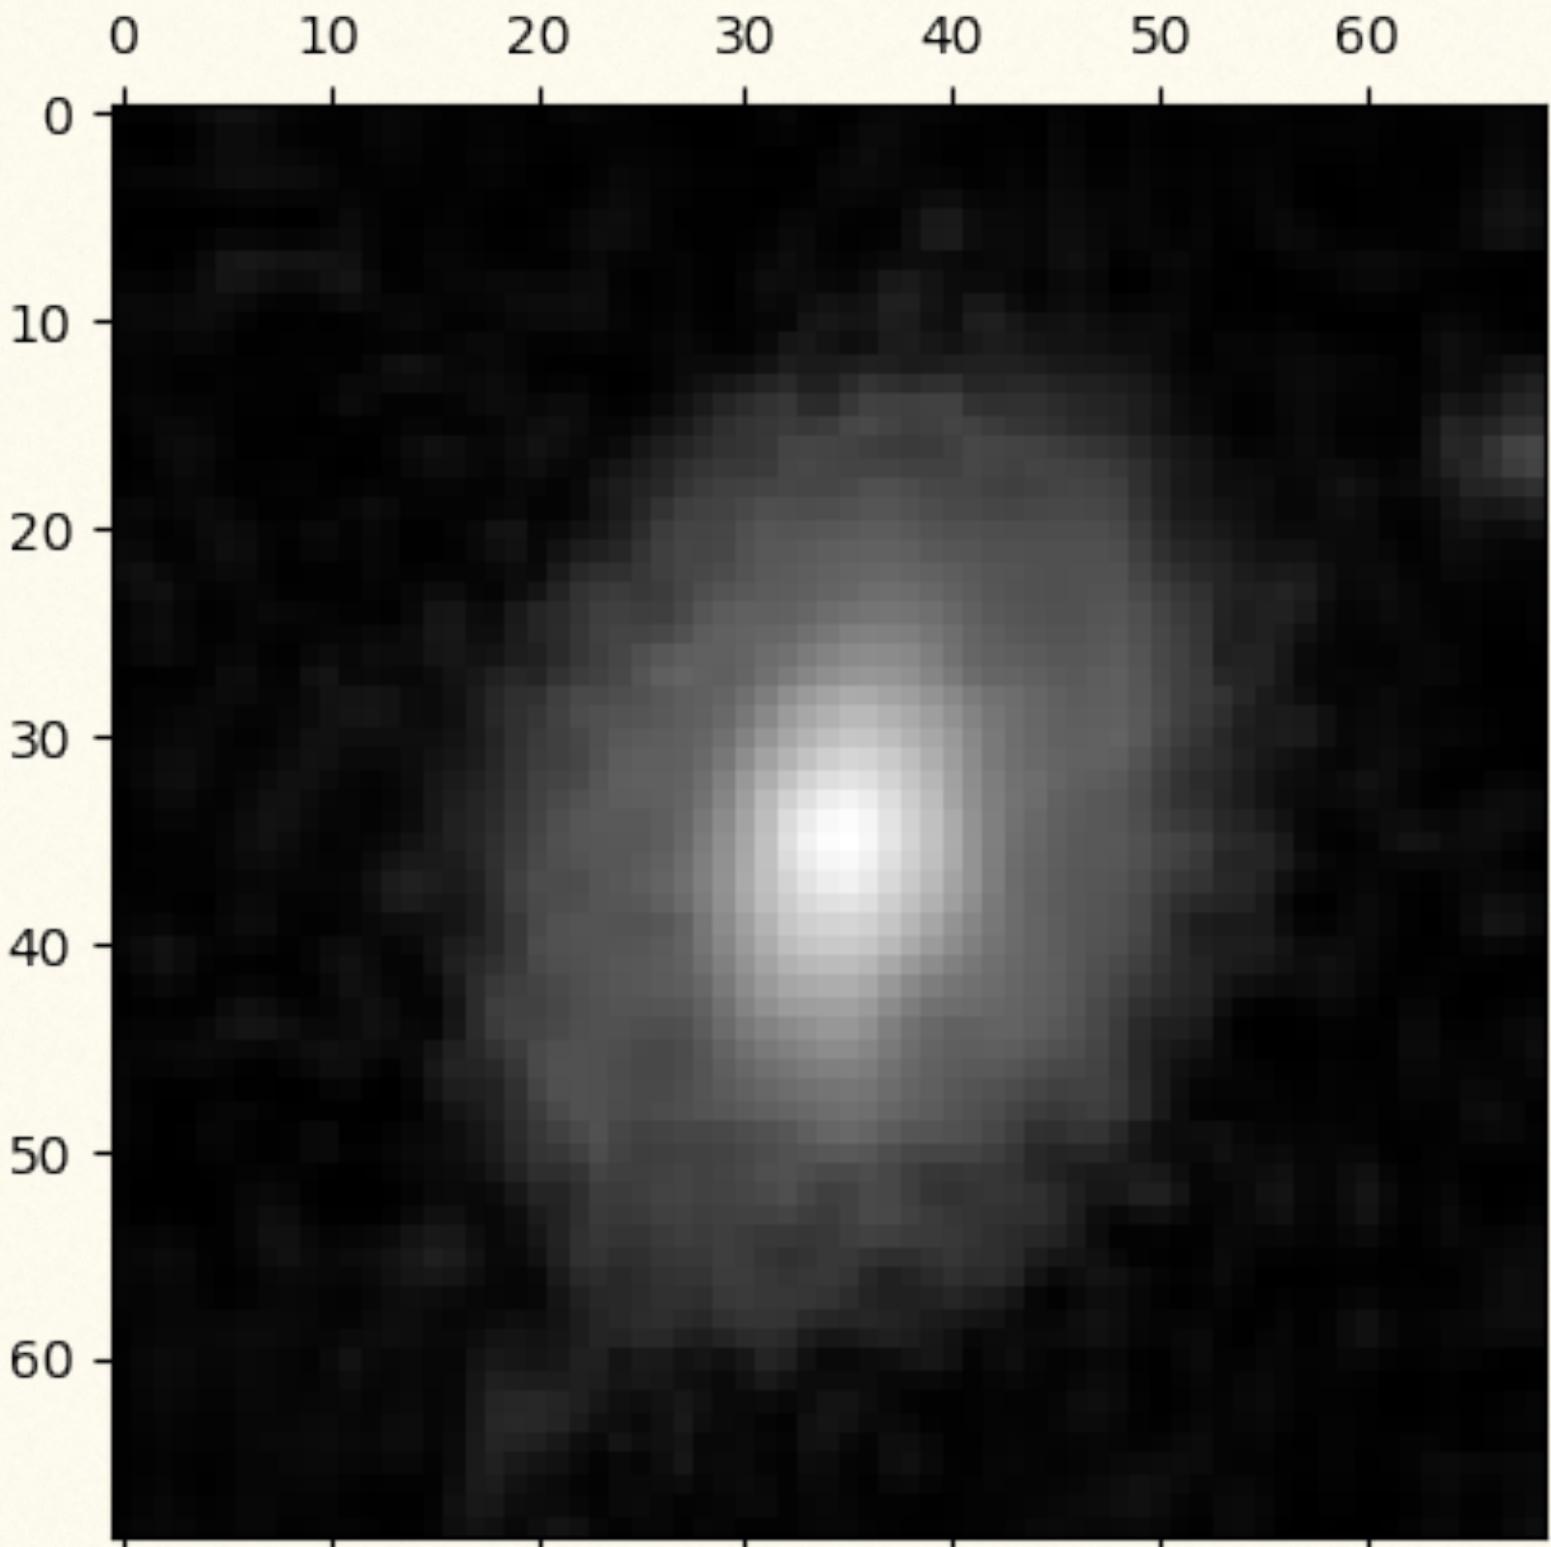
The problem: classifying galaxies

Example images of each class from Galaxy10 dataset



We have down-selected the database to only contain these 5. Representation:
[1,0,0,0,0], [0,1,0,0,0]

Our task



69x69 values with value
[0, 1]

$$f: \mathbb{R}^{4761} \rightarrow \mathbb{R}^5$$

[0.70281476,
0.62053204,
0.7327105,
0.3870779,
0.46099216]

5 numbers that represent
confidence in each
category.

Discussion: How would you parameterize a linear (technically affine) function from $\mathbb{R}^{4761} \rightarrow \mathbb{R}^5$? (This will be a building block!)

Answer:

$$y = Mx + b$$

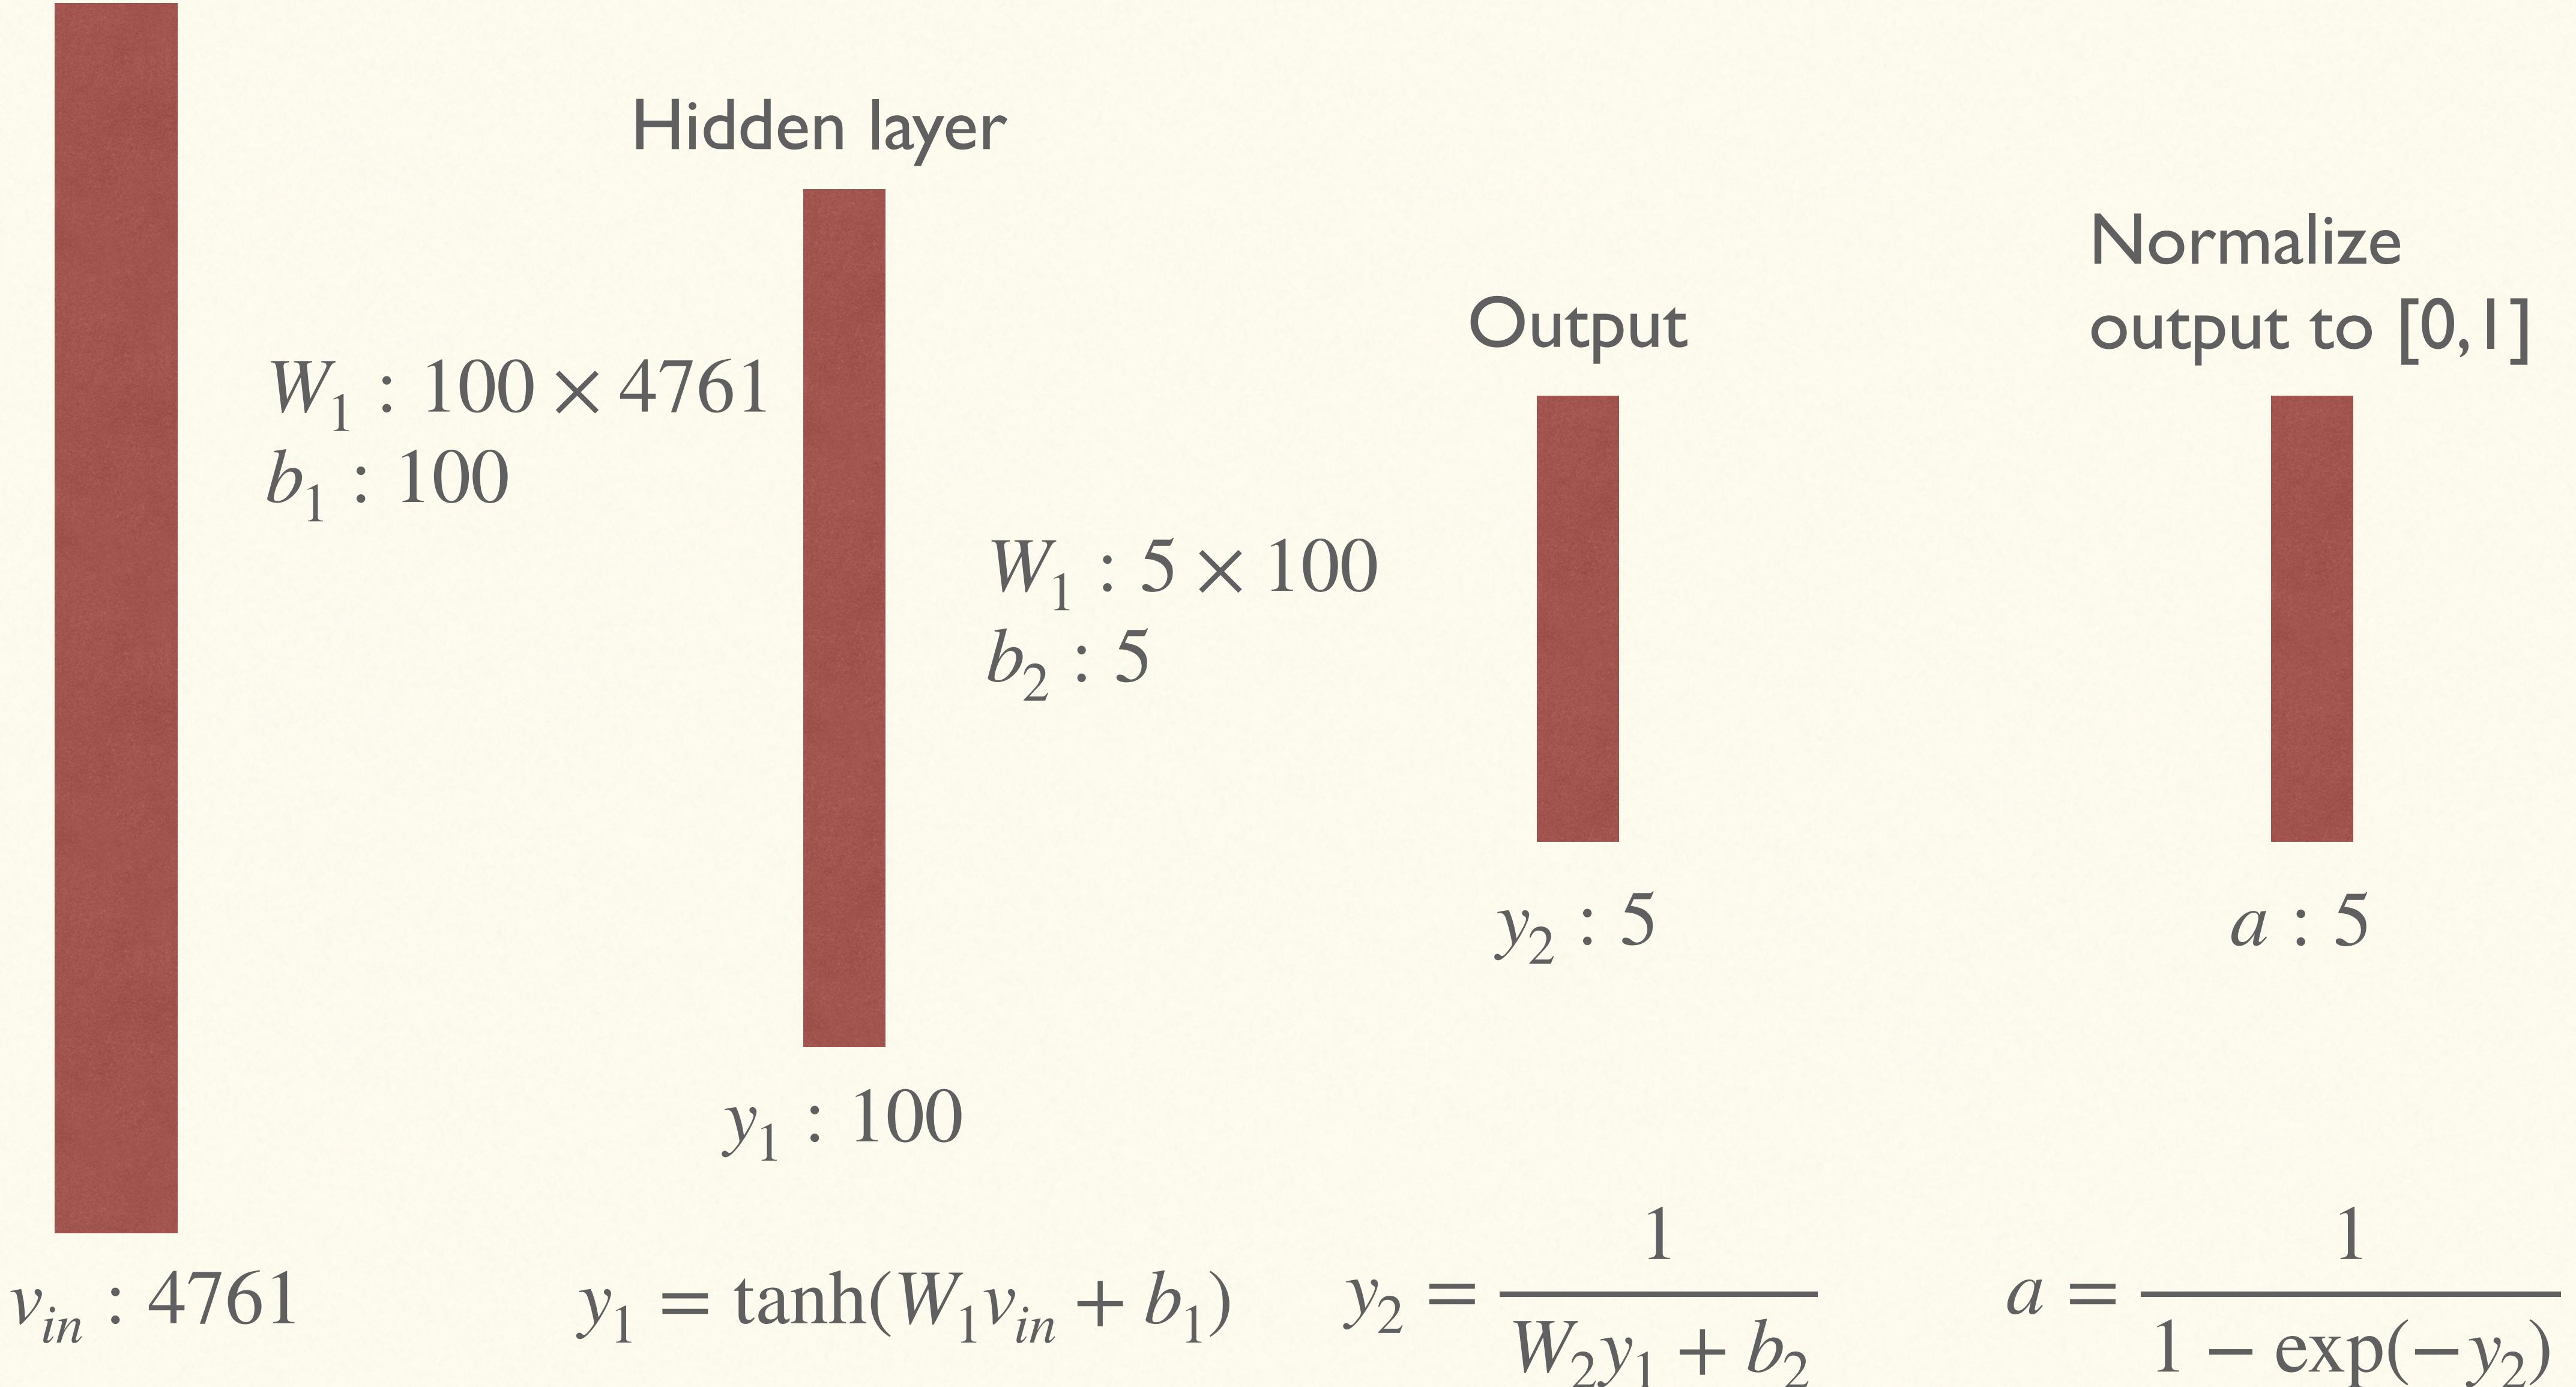
where

$$y, b \in \mathbb{R}^5$$

$$x \in \mathbb{R}^{4761}$$

M is a 5×4761 matrix

A simple neural network: our $f: \mathbb{R}^{4761} \rightarrow \mathbb{R}^5$



Accuracy

Prediction: $a = [0.70281476, 0.62053204, \mathbf{0.7327105}, 0.3870779, 0.46099216]$

Truth: $y = [1. 0. 0. 0. 0.]$

Highest probability is i=2, which is wrong! Return false.

We want to train our model to have the highest accuracy.

Loss function

Prediction: $a = [0.70281476, 0.62053204, 0.7327105, 0.3870779, 0.46099216]$

Truth: $y = [1. 0. 0. 0. 0.]$

Want to have a small loss function when the first number is large:

$$-\log(a_0)$$

And a large loss function when the other numbers are large:

$$\sum_{i>0} \log(1 - a_i)$$

General formula:

$$\sum_i -y_i \log(a_i) - (1 - y_i) \log(1 - a_i) : \text{"cross-entropy"}$$

Discussion:
Why not just the
number of correct
predictions?

Optimizing the loss function

476705 total parameters.

$W_1 : 100 \times 4761$

$b_1 : 100$

$W_1 : 5 \times 100$

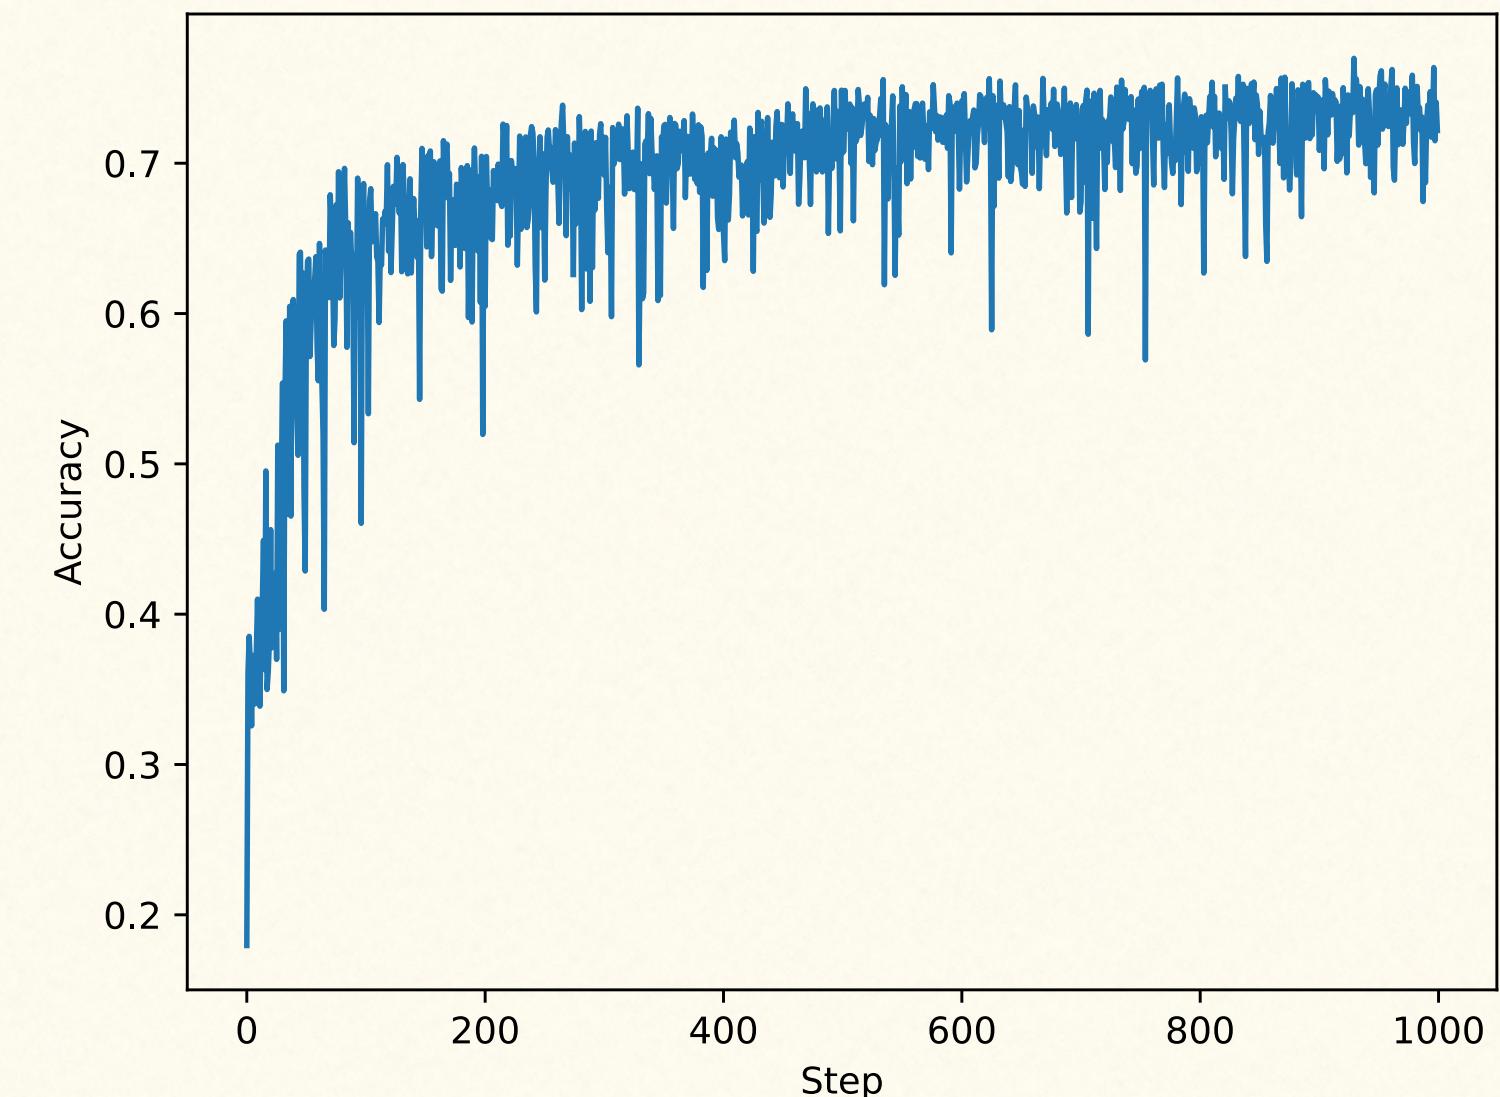
$b_2 : 5$

Second derivative matrix would be 1.7 terabytes of memory.

Stochastic gradient descent:

1. Evaluate the gradient of the loss for ONE random image in the training set
2. Make a tiny change in parameters in the direction of the gradient.
3. Repeat many many times

The reason this works is pretty subtle. It's important that we do not optimize on all images in the set at once; otherwise we will overfit.



Test versus training set: how do you know if we overfit?

Why are test questions different from homework questions?

The test set is reserved and not looked at during training. Your model's performance on the test set will tell you whether

Using jax

Very similar to numpy, but you use `jnp` instead of `np`.

You get:

- 1) just-in-time compilation (speed)
- 2) automatic execution on GPU
- 3) automatic differentiation
- 4) `vmap`

Requires a functional style
(no global variables, no side effects, can't modify arguments)

```
function = jit(function)  
  
cpu=True  
if cpu:  
    jax.config.update('jax_platform_name','cpu')  
    # jax.config.update("jax_enable_x64",True)  
else:  
    pass
```

```
gradient = jax.grad(function)
```

Automatic differentiation (forward mode)

```
#include <iostream>
struct Dual {
    float realPart, infinitesimalPart;
    Dual(float realPart, float infinitesimalPart=0): realPart(realPart),
    infinitesimalPart(infinitesimalPart) {}
    Dual operator+(Dual other) {
        return Dual(
            realPart + other.realPart,
            infinitesimalPart + other.infinitesimalPart
        );
    }
    Dual operator*(Dual other) {
        return Dual(
            realPart * other.realPart,
            other.realPart * infinitesimalPart + realPart * other.infinitesimalPart
        );
    }
}
// Example: Finding the partials of z = x * (x + y) + y * y at (x, y) = (2, 3)
Dual f(Dual x, Dual y) { return x * (x + y) + y * y; }
int main () {
    Dual x = Dual(2);
    Dual y = Dual(3);
    Dual epsilon = Dual(0, 1);
    Dual a = f(x + epsilon, y);
    Dual b = f(x, y + epsilon);
    std::cout << "∂z/∂x = " << a.infinitesimalPart << ", "
          << "∂z/∂y = " << b.infinitesimalPart << std::endl;
    // Output: ∂z/∂x = 7, ∂z/∂y = 8
    return 0;
}
```

Suppose

$$g(x, y) = x + y$$
$$f(x, y) = x \cdot y$$

Then

$$h(x, y) = x \cdot (x + y) + y \cdot y = f(x, g(x, y)) + f(x, y)$$

vmap

Allows you to add dimensions to your arrays.

say:

```
def f(a, x):  
    return a@x
```

takes in a (matrix) and x (vector) and returns a vector

```
fvmap = jax.vmap(f, in_axes=(None, 0), out_axes=0))
```

takes in a (matrix) and x (N vectors) and returns N vectors